

Data mining



What do you do with LOTS of data?

General guidelines:

1. *Look* at your data

- perform any/all preprocessing steps

- use many visualizations, especially of raw, unaveraged data

 - scatterplots, heatplots, histograms

- look at distributions, identify outliers

- test hypotheses with encoding & decoding approaches

 - (confusion matrices!)

- look for trends and patterns (*unsupervised* methods)

 - dimensionality reduction, clustering

What do you do with LOTS of data?

General guidelines:

1. *Look* at your data
2. But don't look at *all* of your data

Inflation bias, a.k.a. “p-hacking”


<u>P-VALUE</u>	<u>INTERPRETATION</u>
0.001	HIGHLY SIGNIFICANT
0.01	
0.02	
0.03	
0.04	SIGNIFICANT
0.049	
0.050	OH CRAP. REDO CALCULATIONS.
0.051	ON THE EDGE OF SIGNIFICANCE
0.06	
0.07	HIGHLY SUGGESTIVE, SIGNIFICANT AT THE $P < 0.10$ LEVEL
0.08	
0.09	
0.099	HEY, LOOK AT THIS INTERESTING SUBGROUP ANALYSIS
≥ 0.1	



What do you do with LOTS of data?

General guidelines:

1. *Look* at your data
2. But don't look at *all* of your data

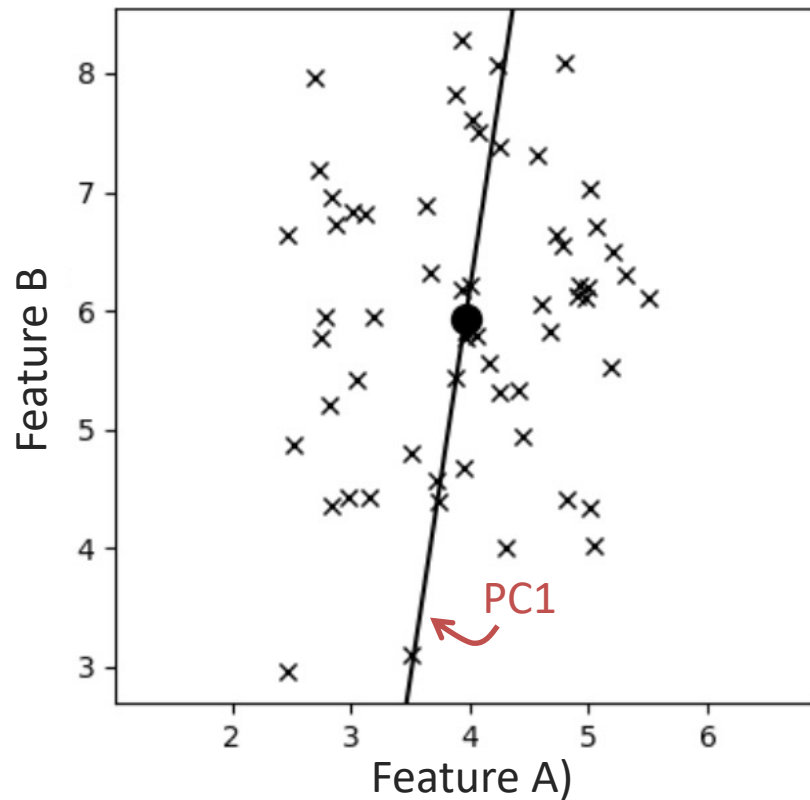


Best practices to avoid p-hacking

1. Create your own replication study by mining half your data and holding half out for test
2. Use visualizations to see patterns, do stats only at the end

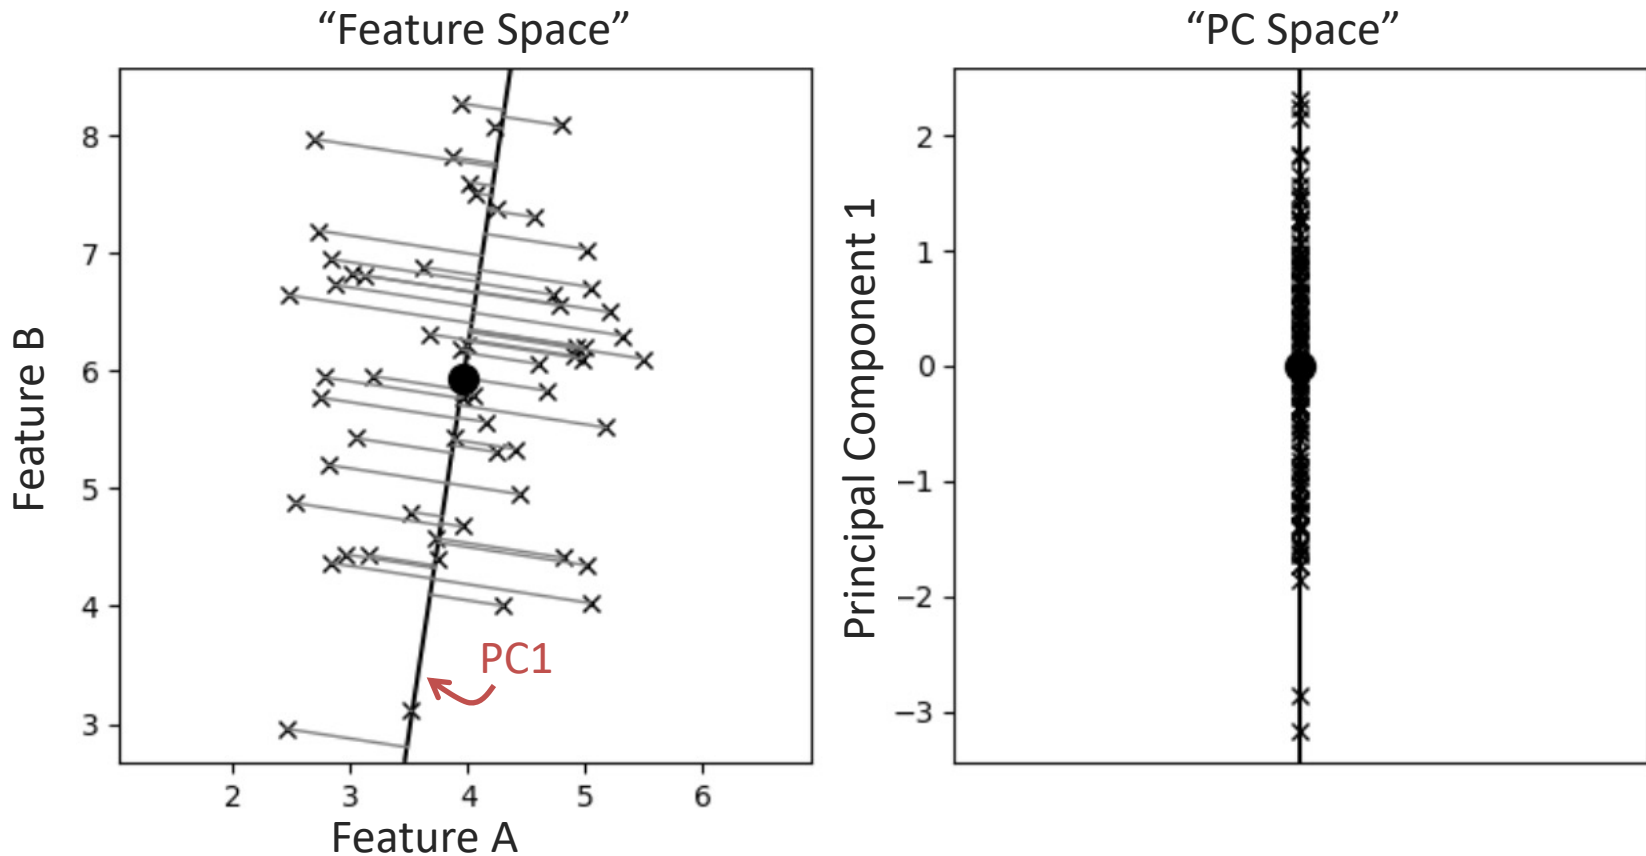
Finding patterns in data: Dimensionality reduction

Principal component analysis (PCA)



Finding patterns in data: Dimensionality reduction

Principal component analysis (PCA)



PCA demo

Let's create some fake data to see how PCA works:

10 features (e.g. neurons)

8 observations (e.g. a measure in 8 different conditions)

All features reflect an underlying pattern:

```
% Let's make some features that follow an underlying patterns
% We'll initialize the pattern as responses observed to 8 different
% conditions (of in 8 different observations)
pattern1 = [5 5 10 10 5 5 10 10];
var = 1;

% Next we'll make a population of 10 features that follow this pattern with
% some noise
pop1 = [];
% let's make 10 features per subpopulation
for k=1:10
    for j=1:8 %there are 8 observations
        noise = normrnd(0,var); %this will add random noise, drawn from a gaussian centered at 0 with standard deviation = var,
        pop1(k,j) = pattern1(j)+noise;
    end
end
```

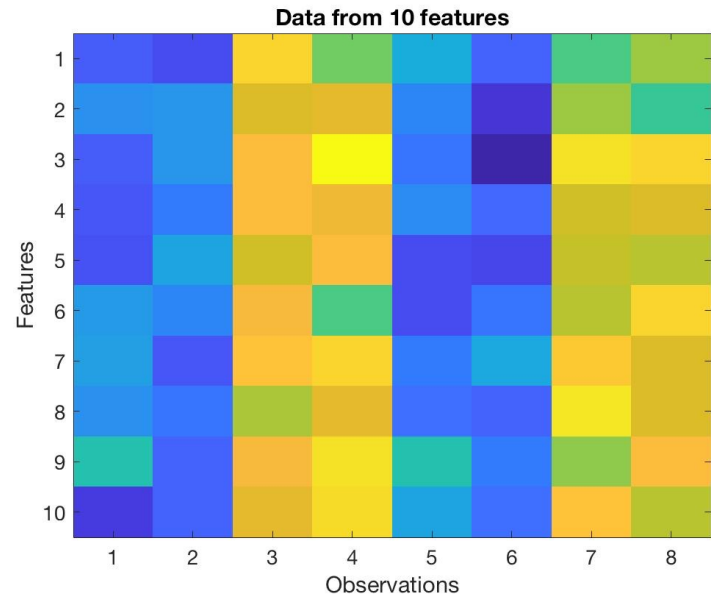

PCA demo

Let's create some fake data to see how PCA works:

10 features (e.g. neurons)

8 observations (e.g. a measure in 8 different conditions)

All features reflect an underlying pattern:



PCA demo

Let's create some fake data to see how PCA works:

10 features (e.g. neurons)

8 observations (e.g. a measure in 8 different conditions)

All features reflect an underlying pattern:

In matlab:

```
% Now let's use PCA to see how it recovers patterns  
[coeff,pcs,~,~,explained]=pca(pop1'); % note that the function that computes PCA has to take input matrices in the correct orientat
```

In R:

```
pop1_pca <- prcomp(pop1, center = FALSE) # note by default matlab "centers" data R does not  
# so scale of values differs between matlab and R
```


PCA demo

Let's create some fake data to see how PCA works:

10 features (e.g. neurons)

8 observations (e.g. a measure in 8 different conditions)

All features reflect an underlying pattern:

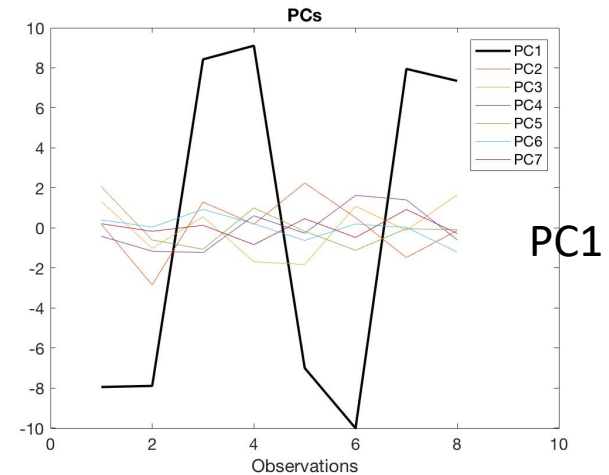


PCA can give you:

- **PCs**

- % variance explained by each PC


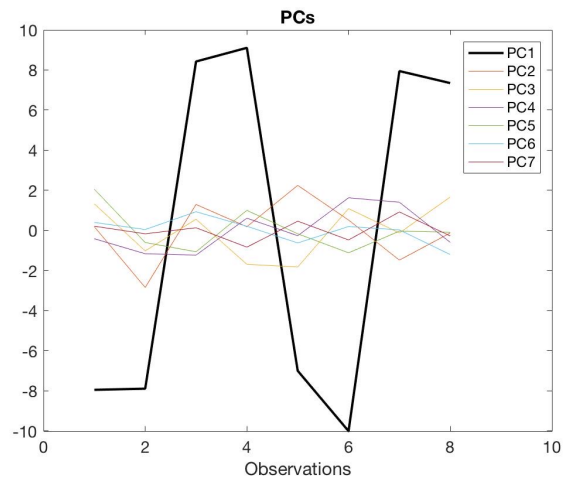
- feature loadings or “weights” for each PC



PCA demo

PCA can give you:

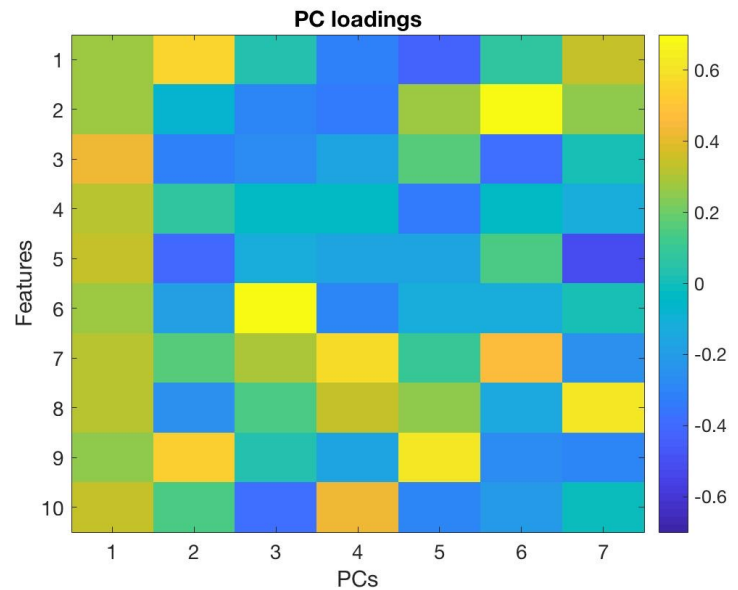

- PCs
- **% variance explained by each PC**
- feature loadings for each PC



PCA demo

PCA can give you:

- PCs
- % variance explained by each PC
- **feature loadings for each PC**



- PC1 has high weights on each feature
- Other PCs weigh on only 1 or 2 features

PCA demo

Let's add data....

30 features (e.g. neurons)

8 observations (e.g. measure in 8 conditions)


3 underlying patterns

```
% Let's do it again with a more complex data set
% Additional patterns
pattern2 = [5 10 5 10 5 10 5 10];
pattern3 = [5 5 5 5 10 10 10 10];


% and create two more populations that follow different patterns
pop2 = []; pop3 = [];
for k = 1:10
    for j = 1:8
        noise = normrnd(0,var); % noise should be independent for this simulation
        pop2(k,j) = pattern2(j)+noise;
        noise = normrnd(0,var);
        pop3(k,j) = pattern3(j) + noise;
    end
end
pop = [pop1;pop2;pop3]; % Our full feature matrix is all of these subpopulations together
```

PCA demo

Let's add data....




30 features (e.g. neurons)
8 observations (e.g. measure in 8 conditions)
3 underlying patterns




PCA demo

Let's add data....


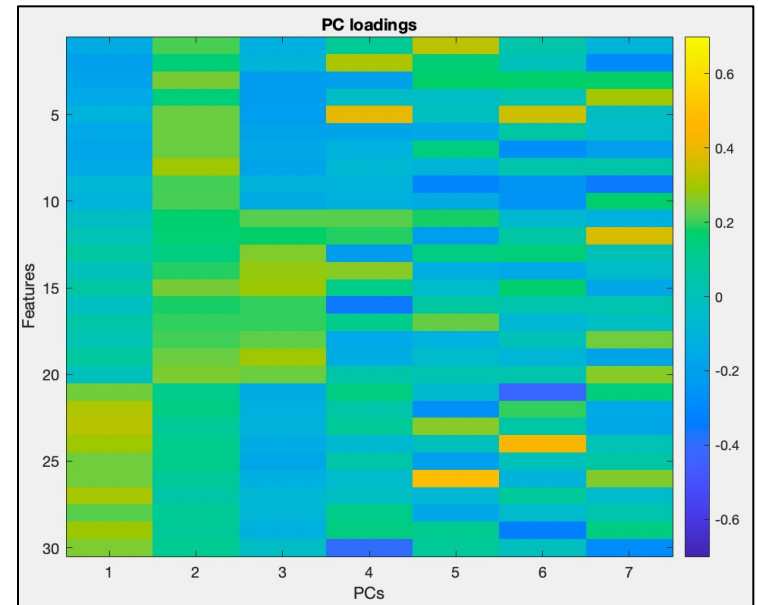



30 features (e.g. neurons)
8 observations (e.g. conditions)
3 underlying patterns



PCA demo

Let's add data....





Finding patterns in data: Dimensionality reduction

Many algorithms

PCA

Singular value decomposition (SVD – closely related to PCA)

Independent component analysis (ICA) - *similar to PCA, but finds components that have maximal statistical independence*



Finding patterns in data: Dimensionality reduction

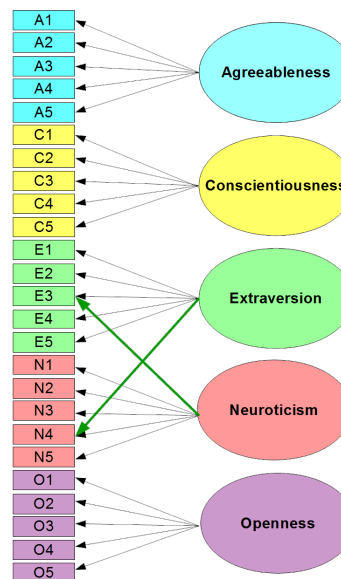
Many algorithms

PCA

Singular value decomposition (SVD – closely related to PCA)

Independent component analysis (ICA)

Factor analysis – *reduces a large number of variables to a smaller number of "factors" or latent variables by finding maximum common variance*



Finding patterns in data: Dimensionality reduction

Many algorithms

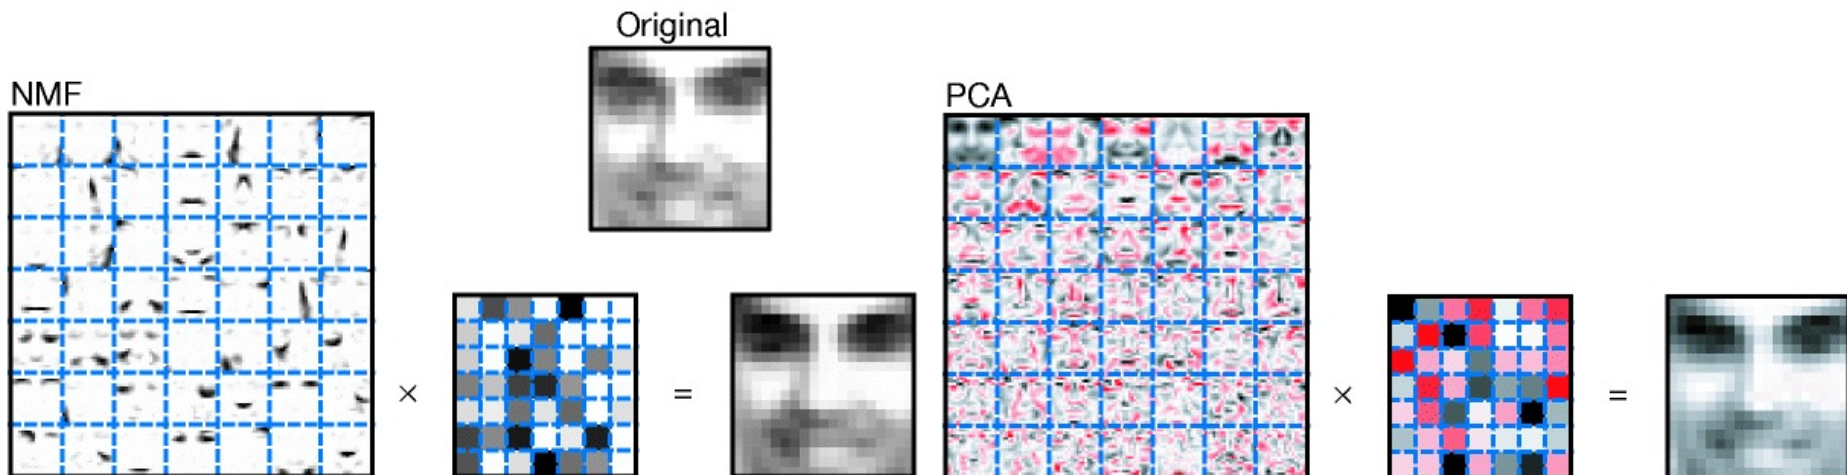
PCA

Singular value decomposition (SVD – closely related to PCA)

Independent component analysis (ICA)

Factor analysis

Non-negative matrix factorization (NMF) – *like PCA but constrains weights to be non-negative -> more interpretable in some instances*



Finding patterns in data: Dimensionality reduction

Many algorithms

PCA

Singular value decomposition (SVD – closely related to PCA)

Independent component analysis (ICA)

Factor analysis

Non-negative matrix factorization (NMF)

demixed PCA (dPCA)


Linear discriminant analysis (LDA)

Bespoke statespace analyses

} *Methods that aim to find interpretable components*


PCA:

component axes that maximize the variance



LDA:

maximizing the component axes for class-separation



Finding patterns in data: Dimensionality reduction

Many algorithms

PCA

Singular value decomposition (SVD – closely related to PCA)

Independent component analysis (ICA)

Factor analysis

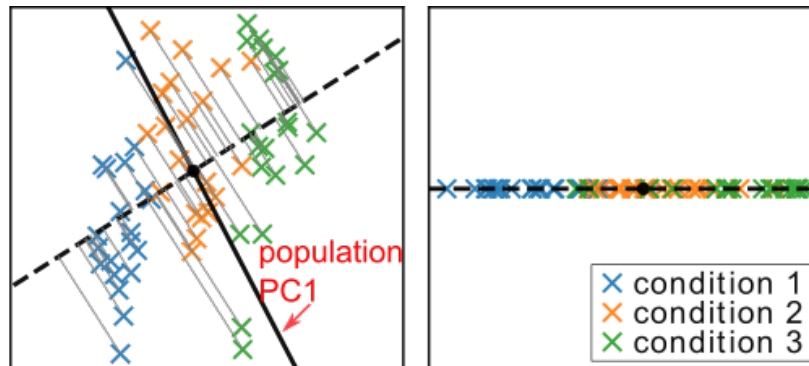
Non-negative matrix factorization (NMF)

demixed PCA (dPCA)

Linear discriminant analysis (LDA)

Bespoke statespace analyses

} *Methods that aim to find interpretable components*



Homework #9

(first part)

HW9: Data mining

You have recorded pupil responses in a subject viewing different images.

The data are saved in *data.txt*, which includes 500 trials where each trial is 1200ms long

1. Plot the mean pupil response over all trials
2. Do PCA across trials (Hint: each PC should be 1200 elements long, and there should be 500 of them).
How much variance does the first PC account for? How many components account for $\geq 90\%$ of the variance?
3. Plot the first principal component.
4. Run k-means 10 times with $k=2$. For each run, plot the cluster centers you obtain.